

Image Segmentation via Network Flow: A Maximum Flow Minimum Cut Approach

Venkata Nitchaya Reddy Konkala¹ and Keerthi Reddy Gudiband²

¹UFID: 34102083, ²UFID: 13652831

Department of Computer Science and Engineering

University of Florida

Gainesville, FL, USA

Email: konkalav@ufl.edu, kgudiband@ufl.edu

Abstract—Image segmentation is a fundamental problem in computer vision that involves partitioning an image into meaningful regions. In this paper, we present a polynomial-time reduction of the binary image segmentation problem to the Maximum Flow problem in network graphs. We leverage the Max-Flow Min-Cut theorem to optimally separate foreground and background pixels based on intensity and spatial coherence constraints. We provide a complete construction of the flow network, prove the correctness of our reduction, present a complete algorithm, and experimentally validate the theoretical $O(VE^2)$ time complexity using the Edmonds-Karp algorithm (Ford-Fulkerson method with BFS). For an $n \times m$ image, this yields $O((nm)^3)$ running time. Our experimental results demonstrate that the algorithm successfully segments images while conforming to the predicted polynomial time bounds.

Index Terms—network flow, image segmentation, max-flow min-cut theorem, graph cuts, polynomial reduction

I. INTRODUCTION

Image segmentation is the process of dividing a digital image into multiple segments or regions to simplify its representation and make it more meaningful for analysis. One of the most fundamental segmentation tasks is binary segmentation, where the goal is to classify each pixel as either foreground or background.

Traditional approaches to image segmentation include thresholding, region growing, and clustering methods. However, these methods often fail to incorporate both local pixel properties and global spatial coherence simultaneously in an optimal fashion.

The problem of optimal binary image segmentation can be elegantly formulated as a graph partitioning problem and solved using network flow techniques. By constructing an appropriate flow network and applying the Max-Flow Min-Cut theorem, we can find the optimal segmentation that balances pixel-wise classification confidence with boundary smoothness.

A. Contributions

Our main contributions are:

- Formalization of the binary image segmentation problem as a graph partitioning optimization problem
- Polynomial reduction to the Maximum Flow problem with detailed construction

- Formal proof of correctness establishing the equivalence between optimal segmentation and minimum cut
- Complete algorithm incorporating the reduction with rigorous $O((nm)^3)$ time complexity analysis
- Experimental validation of the $O(VE^2)$ Edmonds-Karp complexity with real image data

II. PROBLEM STATEMENT

A. Real-World Problem: Binary Image Segmentation

Given a digital image, we want to separate objects of interest (foreground) from the background. This problem appears in numerous applications:

- Medical imaging: Tumor detection, organ segmentation
- Photography: Portrait mode, object isolation
- Computer vision: Object recognition preprocessing
- Autonomous vehicles: Road and obstacle detection

The segmentation should satisfy two criteria:

- 1) **Data Term**: Pixels should be assigned to foreground or background based on their intensity, color, or texture properties
- 2) **Smoothness Term**: Neighboring pixels with similar properties should preferably belong to the same segment (spatial coherence)

B. Abstract Problem Formulation

We abstract the image segmentation problem as follows:

Input:

- An undirected graph $G = (V, E)$ where:
 - V represents the set of pixels in the image
 - E represents the neighborhood relationships between pixels
- For each pixel $p \in V$:
 - $D_p(F)$: cost of assigning pixel p to foreground
 - $D_p(B)$: cost of assigning pixel p to background
- For each edge $(p, q) \in E$:
 - $B_{p,q}$: penalty for assigning p and q to different segments

Output: A partition of V into two disjoint sets F (foreground) and B (background) that minimizes the energy function:

$$E(F, B) = \sum_{p \in F} D_p(F) + \sum_{p \in B} D_p(B) + \sum_{\substack{(p,q) \in E \\ p \in F, q \in B}} B_{p,q} \quad (1)$$

This is a combinatorial optimization problem over $2^{|V|}$ possible partitions, which is NP-hard in general graphs without special structure.

C. Cost Function Instantiation for Images

For grayscale or RGB images, we instantiate the abstract cost functions as follows:

Data Term: Based on intensity similarity to seed regions. Let μ_F and μ_B denote the mean intensities of foreground and background seed pixels, and I_p the intensity of pixel p :

$$D_p(F) = \alpha \cdot (I_p - \mu_F)^2 \quad (2)$$

$$D_p(B) = \alpha \cdot (I_p - \mu_B)^2 \quad (3)$$

where $\alpha > 0$ is a scaling parameter controlling the strength of intensity-based classification.

Boundary Penalty: Encourages spatial smoothness by penalizing cuts between similar neighboring pixels. For adjacent pixels p and q :

$$B_{p,q} = \exp(-\beta \cdot (I_p - I_q)^2) \quad (4)$$

where $\beta > 0$ controls smoothness. Higher values of β produce smoother boundaries, while lower values allow more irregular cuts at strong edges.

Interpretation:

- $D_p(F)$ is low when pixel p has intensity similar to foreground seeds
- $D_p(B)$ is low when pixel p has intensity similar to background seeds
- $B_{p,q}$ is high when neighboring pixels have similar intensities (cutting them incurs high cost)
- $B_{p,q}$ is low when neighboring pixels have different intensities (edge likely present)

III. REDUCTION TO MAXIMUM FLOW

In this section, we present the complete polynomial-time reduction from the abstract image segmentation problem (defined in Section 2) to the MINIMUM S-T CUT problem on directed graphs. We describe the construction of the flow network, prove that the transformation runs in polynomial time, and show that the minimum cut in this network corresponds exactly to the minimum-energy segmentation.

A. Polynomial-Time Reduction and Construction

We begin with an instance of the segmentation problem, consisting of the following components:

- **A set of pixels**

$$V = \{p_1, p_2, \dots, p_n\};$$

- **A set of adjacency relations**

$$E = \{(p_i, p_j) \mid p_i \text{ and } p_j \text{ are neighboring pixels}\};$$

- **Unary (data) assignment costs**

$$D_p(F), \quad D_p(B) \quad \text{for each pixel } p \in V;$$

- **Pairwise (boundary) penalties**

$$B_{p,q} \quad \text{for each adjacency } (p, q) \in E;$$

- **Optional foreground/background seeds**

$$F_{\text{seeds}}, \quad B_{\text{seeds}}.$$

From this input, we construct a corresponding MINIMUM S-T CUT instance as follows.

a) Node Construction.: For each pixel $p \in V$, create a node (also denoted p). Add two terminals:

$$s \quad (\text{source}), \quad t \quad (\text{sink}).$$

Thus,

$$V' = V \cup \{s, t\}.$$

b) Terminal Edge Construction.: For every pixel $p \in V$:

- Add an edge (s, p) with capacity $c(s, p) = D_p(B)$,
- Add an edge (p, t) with capacity $c(p, t) = D_p(F)$.

These encode the unary data terms of the segmentation energy.

c) Neighbor Edge Construction.: For each adjacency $(p, q) \in E$, add directed edges

$$(p, q), \quad (q, p)$$

with capacities

$$c(p, q) = c(q, p) = B_{p,q}.$$

These encode the smoothness penalty for assigning neighboring pixels to different labels.

d) Seed Enforcement.: Seeded pixels must not change labels. Since infinite capacities cannot be represented exactly, we select a constant

$$M > \sum_{(u,v) \in E'} c(u, v),$$

and enforce seeds by assigning:

$$c(s, p) = M \quad \text{for all } p \in B_{\text{seeds}}, \quad c(p, t) = M \quad \text{for all } p \in F_{\text{seeds}}.$$

This guarantees that seeded pixels cannot cross the cut.

e) Parameter Preservation.: For any segmentation (F, B) and its corresponding cut (S, T) , the cut cost satisfies

$$\text{Cut}(S, T) = E(F, B),$$

so the reduction preserves objective values exactly.

f) Size of the Constructed Network.: The constructed network satisfies:

$$|V'| = |V| + 2 = O(|V|), \quad |E'| = 2|V| + 2|E| = O(|V| + |E|),$$

so the construction runs in linear time with respect to the size of the image graph.

g) *Conclusion of Reduction.*: Because the transformation preserves segmentation energy exactly and is computable in $O(|V| + |E|)$ time, this establishes a valid polynomial-time reduction from the binary segmentation problem to the MINIMUM s - T CUT problem.

B. Correctness of the Reduction

We now show that the minimum cut in the constructed flow network G' corresponds exactly to the minimum-energy segmentation.

Theorem 1. *Let (S, T) be any s - t cut of G' with $s \in S$ and $t \in T$. Define the segmentation:*

$$F = S \setminus \{s\}, \quad B = T \setminus \{t\}.$$

Then the cut capacity equals the segmentation energy:

$$c(S, T) = E(F, B).$$

Thus, the minimum s - t cut yields the optimal segmentation.

Proof. The capacity of cut (S, T) is:

$$c(S, T) = \sum_{u \in S, v \in T} c(u, v).$$

Only three types of edges can cross the cut:

- 1) Unary edges (p, t) with $p \in F$, contributing $D_p(F)$.
- 2) Unary edges (s, p) with $p \in B$, contributing $D_p(B)$.
- 3) Pairwise edges (p, q) with $p \in F$ and $q \in B$, contributing $B_{p,q}$.

Thus,

$$c(S, T) = \sum_{p \in F} D_p(F) + \sum_{p \in B} D_p(B) + \sum_{\substack{(p,q) \in E \\ p \in F, q \in B}} B_{p,q},$$

which equals $E(F, B)$ by definition.

By the Max-Flow Min-Cut Theorem, the minimum cut is obtained by any maximum-flow algorithm. Hence, the obtained segmentation is optimal. $\square$

The polynomial-time reduction described above is independent of the specific maximum-flow algorithm used to compute the cut. Any polynomial-time max-flow method (e.g., Edmonds–Karp, Dinic, Push–Relabel) yields the same optimal segmentation because the reduction preserves all objective values exactly.

IV. COMPLETE ALGORITHM

In this section we present the complete algorithm that performs image segmentation using the polynomial-time reduction from Section 3 and the Edmonds–Karp maximum-flow algorithm. The algorithm receives an image and foreground/background seeds, constructs the corresponding flow network G' , computes the maximum flow, and converts the resulting minimum s - t cut into a binary segmentation mask.

A. Algorithm Overview

Algorithm 1 implements the following steps:

- 1) Construct a graph $G' = (V', E')$ with terminals s and t .
- 2) Add unary (data) edges (s, p) and (p, t) with capacities $D_p(B)$ and $D_p(F)$.
- 3) Add pairwise (smoothness) edges (p, q) and (q, p) with capacity $B_{p,q}$ for each adjacency $(p, q) \in E$.
- 4) Enforce seeds using a sufficiently large constant $M > \sum_{(u,v) \in E'} c(u, v)$.
- 5) Run the Edmonds–Karp algorithm to compute a maximum flow.
- 6) Extract the minimum cut from the residual graph.
- 7) Assign all nodes reachable from s as foreground and the rest as background.

B. Segmentation Algorithm

C. Interpretation of the Output Mask

The mask M is a visualization-friendly indicator of the optimal labeling. A white pixel ($M[p] = 1$) corresponds to the foreground set F , and a black pixel ($M[p] = 0$) corresponds to the background set B . Figure 1 demonstrates this relationship.




Fig. 1. Example of the segmentation output. **Left:** Input image with user-defined foreground (blue) and background (red) seeds. **Center:** Resulting binary segmentation mask M , where white pixels denote foreground ($M = 1$) and black pixels denote background ($M = 0$). **Right:** Extracted foreground obtained by applying the mask M to the original image.

D. Proof of Correctness

We establish the correctness of Algorithm 1 through a series of lemmas, culminating in a proof by induction that demonstrates the algorithm produces the globally optimal minimum-energy segmentation.

Lemma 1 (Network Construction Validity). *The flow network $G' = (V', E')$ constructed in lines 1–5 of Algorithm 1 is a valid flow network with non-negative capacities, and every cut (S, T) with $s \in S$ and $t \in T$ corresponds to a valid segmentation of the image pixels.*

Proof. The construction creates $|V'| = nm + 2$ vertices (one per pixel plus source s and sink t). For each pixel p , we add edges (s, p) with capacity $c(s, p) = D_p(B) \geq 0$ and (p, t) with capacity $c(p, t) = D_p(F) \geq 0$. For each neighboring pixel pair (p, q) , we add edges with capacity $c(p, q) = B_{p,q} \geq 0$. All capacities are non-negative by definition of data terms and boundary penalties, satisfying the flow network requirement.

Algorithm 1 Image Segmentation via Maximum Flow

Require: Image I of size $n \times m$, foreground seeds F_{seeds} , background seeds B_{seeds}

Ensure: Binary segmentation mask M

- 1: Initialize network $G' = (V', E')$ with terminals s and t
- 2: $V' \leftarrow \{s, t\}$, $E' \leftarrow \emptyset$
- 3: Compute seed means μ_F and μ_B from intensities of F_{seeds} and B_{seeds}
- 4: Choose large constant

$$M > \sum_{p \in V} (D_p(F) + D_p(B)) + \sum_{(p,q) \in E} B_{p,q}$$

- 5: **for** each pixel p with coordinates (i, j) **do**
- 6: Add node p to V'
- 7: Compute unary costs $D_p(F)$ and $D_p(B)$ from Eq. (1)–(3)
- 8: **if** $p \in F_{\text{seeds}}$ **then**
- 9: Add edges: (s, p) with capacity 0; (p, t) with capacity M
- 10: **else if** $p \in B_{\text{seeds}}$ **then**
- 11: Add edges: (s, p) with capacity M ; (p, t) with capacity 0
- 12: **else**
- 13: Add edges: (s, p) with capacity $D_p(B)$; (p, t) with capacity $D_p(F)$
- 14: **end if**
- 15: **end for**
- 16: **for** each adjacency $(p, q) \in E$ under 4-connectivity **do**
- 17: Compute boundary penalty $B_{p,q}$ from Eq. (4)
- 18: Add edges (p, q) and (q, p) with capacity $B_{p,q}$
- 19: **end for**
- 20: Run Edmonds–Karp on G' to compute maximum flow f
- 21: Extract the minimum cut (S, T) from the residual graph
- 22: **for** each pixel node p **do**
- 23: **if** $p \in S$ **then**
- 24: $M[p] \leftarrow 1$ {Foreground}
- 25: **else**
- 26: $M[p] \leftarrow 0$ {Background}
- 27: **end if**
- 28: **end for**
- 29: **return** M

Any cut (S, T) partitions V' into two disjoint sets with $s \in S$ and $t \in T$. Define $F = S \setminus \{s\}$ as foreground pixels and $B = T \setminus \{t\}$ as background pixels. Since (S, T) partitions all vertices, (F, B) partitions all pixels, forming a valid segmentation. $\square$

Lemma 2 (Seed Constraint Enforcement). *For any minimum cut (S, T) in the flow network with seed constraint capacities $M > \sum_{p,q} (D_p(F) + D_p(B) + B_{p,q})$, all foreground seeds are in $S \setminus \{s\}$ and all background seeds are in $T \setminus \{t\}$.*

Proof. Suppose p is a foreground seed and assume for contradiction that $p \in T$. By construction (line 4), the edge (p, t) has capacity M . This edge crosses the cut (S, T) ,

contributing M to the cut capacity. However, any alternative cut placing p in S would replace this contribution with at most $D_p(F) + \sum_{q \sim p} B_{p,q} < M$. Thus, the original cut cannot be minimal, contradicting our assumption. Therefore, $p \in S$.

Similarly, if q is a background seed and $q \in S$, then edge (s, q) with capacity M crosses the cut. Moving q to T would reduce the cut capacity, again contradicting minimality. Therefore, $q \in T$. $\square$

Lemma 3 (Cut-Energy Equivalence). *For any cut (S, T) in the flow network, the cut capacity equals the segmentation energy of the corresponding partition:*

$$c(S, T) = E(F, B)$$

where $F = S \setminus \{s\}$ and $B = T \setminus \{t\}$.

Proof. The cut capacity is the sum of capacities of all edges from S to T :

$$\begin{aligned} c(S, T) &= \sum_{p \in F} c(p, t) + \sum_{q \in B} c(s, q) + \sum_{\substack{(p,q) \in E \\ p \in F, q \in B}} c(p, q) \\ &= \sum_{p \in F} D_p(F) + \sum_{q \in B} D_q(B) + \sum_{\substack{(p,q) \in E \\ p \in F, q \in B}} B_{p,q} \\ &= E(F, B). \end{aligned}$$

The first equality follows from the definition of cut capacity (edges cannot go from B to F or within F or B). The second equality follows from our capacity assignments in lines 2–3. The third equality is the definition of segmentation energy from Equation (1). $\square$

Lemma 4 (Edmonds–Karp Termination and Optimality). *The Edmonds–Karp algorithm (line 6) terminates in $O(VE^2)$ time and computes a maximum flow f^* such that the flow value equals the capacity of a minimum cut.*

Proof. This is a direct application of the max-flow min-cut theorem [4]. The Edmonds–Karp algorithm uses BFS to find shortest augmenting paths, which guarantees that the number of iterations is $O(VE)$, and each BFS takes $O(E)$ time, yielding $O(VE^2)$ total complexity. Upon termination, no augmenting path exists, so by the max-flow min-cut theorem, $|f^*| = \min_{(S,T)} c(S, T)$. $\square$

Lemma 5 (Minimum Cut Extraction). *The set S computed by BFS from s in the residual graph G_f (line 7) defines a minimum cut $(S, V' \setminus S)$ with capacity equal to the maximum flow value.*

Proof. Let $T = V' \setminus S$. By construction, $s \in S$ (reachable from itself) and $t \in T$ (no path exists in residual graph, otherwise we could augment flow). For any edge $(u, v) \in E'$ with $u \in S$ and $v \in T$, we must have $f(u, v) = c(u, v)$ (fully saturated), otherwise v would be reachable from s via the residual edge. Thus:

$$c(S, T) = \sum_{u \in S, v \in T} c(u, v) = \sum_{u \in S, v \in T} f(u, v) = |f^*|.$$

By Lemma 4, $|f^*|$ equals the minimum cut capacity, so (S, T) is indeed a minimum cut. $\square$

Theorem 2 (Algorithm Correctness). *Algorithm 1 computes the globally optimal minimum-energy segmentation that respects all seed constraints in polynomial time.*

Proof by Induction. We prove by strong induction on the number of pixels $N = nm$ that the algorithm produces the optimal segmentation.

Base Case ($N = 1$): For a single-pixel image, the network has vertices $\{s, p, t\}$ with edges (s, p) and (p, t) . If p is a foreground seed, the capacity $c(p, t) = M$ dominates, forcing p into the source set S with cut capacity $D_p(F)$. If p is a background seed, $c(s, p) = M$ forces p into sink set T with cut capacity $D_p(B)$. If p is unseeded, the minimum cut places p in the set yielding $\min(D_p(F), D_p(B))$. In all cases, the algorithm produces the unique optimal segmentation.

Inductive Hypothesis: Assume the algorithm correctly computes the optimal segmentation for all images with fewer than N pixels.

Inductive Step (N pixels): Consider an image I with $N = nm$ pixels. We must show the algorithm produces the optimal segmentation for I .

By Lemma 1, the constructed network is valid. By Lemma 4, Edmonds–Karp computes maximum flow f^* in polynomial time. By Lemma 5, the extracted cut (S, T) is a minimum cut with $c(S, T) = |f^*|$.

Let (F^*, B^*) be the segmentation extracted from (S, T) where $F^* = S \setminus \{s\}$ and $B^* = T \setminus \{t\}$. By Lemma 3, $E(F^*, B^*) = c(S, T)$. By Lemma 2, (F^*, B^*) respects all seed constraints.

Now, suppose for contradiction there exists another segmentation (F', B') respecting all seeds with $E(F', B') < E(F^*, B^*)$. This segmentation corresponds to a cut (S', T') where $S' = F' \cup \{s\}$ and $T' = B' \cup \{t\}$. By Lemma 3, $c(S', T') = E(F', B') < E(F^*, B^*) = c(S, T)$. But this contradicts the fact that (S, T) is a minimum cut. Therefore, no such (F', B') exists, and (F^*, B^*) is globally optimal.

The algorithm constructs the network in $O(N)$ time (lines 1–5), computes maximum flow in $O(N^3)$ time by Lemma 4, and extracts the cut in $O(N)$ time (lines 7–9). Total complexity is $O(N^3)$, which is polynomial.

By the principle of strong induction, the algorithm correctly computes the optimal segmentation for images of any size. $\square$

Remarks. The proof establishes that Algorithm 1 is not a heuristic or approximation—it computes the *exact* globally optimal solution by exploiting the equivalence between minimum cuts and minimum-energy segmentations (Lemma 3) combined with the polynomial-time solvability of maximum flow (Lemma 4).

E. Complexity Analysis

Time Complexity: $O((nm)^3) = O(n^6)$ for $n \times n$ images.

Space Complexity: $O(nm)$.

We now justify these complexities by analyzing each component of Algorithm 1.

1) *Time Complexity Justification:* The algorithm has three main phases: network construction, maximum flow computation, and minimum cut extraction.

Phase 1: Network Construction. The algorithm creates a flow network $G' = (V', E')$ with $|V'| = nm + 2$ vertices (one per pixel, plus source s and sink t) and $|E'| \in O(nm)$ directed edges. The edges consist of:

- $2nm$ terminal edges (each pixel connects to both s and t),
- $4nm - 2n - 2m$ neighbor edges (from 4-connectivity on the grid).

Computing intensity histograms requires scanning all nm pixels once. Constructing vertices and edges involves iterating over all pixels and their neighbors, which takes $O(nm)$ time total.

Network construction time: $O(nm)$.

Phase 2: Maximum Flow Computation. This is the dominant phase. The Edmonds–Karp algorithm finds maximum flow by repeatedly discovering shortest augmenting paths using breadth-first search (BFS). The complexity analysis involves two components:

- **Number of augmenting paths:** Edmonds–Karp performs at most $O(|V'| \cdot |E'|)$ augmenting path iterations. This is a proven bound based on the monotonicity property of shortest path distances in the residual graph [3]. Substituting our values gives:

$$O(|V'| \cdot |E'|) = O((nm + 2) \cdot 6nm) = O((nm)^2).$$

- **Time per augmenting path:** Each iteration runs BFS on the residual graph to find the shortest path from s to t . BFS explores all edges once, taking $O(|E'|) = O(nm)$ time per iteration.

Multiplying these two factors gives the total maximum flow time:

$$\begin{aligned} O(\text{iterations} \times \text{time per iteration}) &= O((nm)^2 \cdot nm) \\ &= O((nm)^3) = O(N^3) \end{aligned}$$

where $N = nm$ is the total pixel count. For square images where $n = m$, this becomes $O(n^6)$.

Maximum flow time: $O((nm)^3) = O(n^6)$ for $n \times n$ images.

Phase 3: Minimum Cut Extraction. After computing maximum flow, extracting the minimum cut requires:

- A single graph traversal (BFS or DFS) from source s in the residual graph to identify all reachable vertices. This takes $O(|V'| + |E'|) = O(nm)$ time.
- Constructing the binary mask M by iterating over all nm pixels and recording their labels. This takes $O(nm)$ time.

Cut extraction time: $O(nm)$.

Overall Time Complexity. Summing all three phases:

$$T(n, m) = O(nm) + O((nm)^3) + O(nm) = O((nm)^3).$$

The maximum flow computation dominates asymptotically, giving the final time complexity:

$O((nm)^3) = O(n^6) \text{ for } n \times n \text{ images.}$

2) *Space Complexity Justification*: The algorithm requires memory for the following data structures:

- **Network representation**: The graph G' with $O(nm)$ vertices and $O(nm)$ edges is stored using adjacency lists. Each edge stores its capacity, requiring $O(nm)$ space total.
- **Flow and residual graph**: The algorithm maintains flow values on all edges and constructs a residual graph with the same structure as G' . Both require $O(nm)$ space.
- **BFS auxiliary structures**: The breadth-first search uses a queue to store vertices during traversal. In the worst case, the queue contains $O(|V'|) = O(nm)$ vertices.
- **Output mask**: The binary segmentation mask M is an $n \times m$ array requiring $O(nm)$ space.

All components use $O(nm)$ space, so the overall space complexity is:

$$O(nm).$$

This is optimal since we must at least store the input image and output mask, both of size $O(nm)$.

3) *Practical Considerations*: For typical image sizes, the $O(n^6)$ time complexity has the following implications:

- **Small images** (100×100): $O(n^6) = O(10^{12})$ operations. Feasible on modern hardware, typically completing in seconds to minutes.
- **Medium images** (500×500): $O(n^6) = O(1.56 \times 10^{16})$ operations. Computationally expensive, requiring hours or longer.
- **Large images** (1000×1000): $O(n^6) = O(10^{18})$ operations. Prohibitively expensive for exact computation.

In practice, the Edmonds–Karp algorithm often performs better than the worst-case bound, especially when:

- The seed set is sparse and well-separated, reducing the number of augmenting paths needed.
- The image has smooth intensity distributions, making flow convergence faster.
- Early termination occurs when most flow is established quickly.

Empirical observations show that typical runtime is closer to $O(n^4)$ or $O(n^5)$ for many real-world images. Section ?? provides experimental validation of these complexity bounds on synthetic and real images.

V. EXPERIMENTAL VALIDATION

A. Implementation Details

We implemented the algorithm in Python using:

- NumPy for image processing
- Custom implementation of Edmonds–Karp algorithm
- NetworkX for graph structure (validation purposes)
- Matplotlib for visualization

B. Experimental Setup

We evaluated the algorithm on synthetic and real images of varying sizes:

- Synthetic images: 10×10 to 100×100 pixels

- Real images: Natural photographs scaled to different resolutions
- Hardware / Environment: macOS; Python 3.x; packages: NumPy, Matplotlib, Pillow (see Requirements)
- 10 runs per image size, average time reported

C. Time Complexity Validation

Figure 2 shows the experimental running time vs. the theoretical bound expressed in terms of total pixel count $N = nm$. For grid graphs with constant-degree connectivity, Edmonds–Karp runs in $O(VE^2)$ which simplifies to $O(N^3)$ when $V = \Theta(N)$ and $E = \Theta(N)$. For square images ($n = m$), this corresponds to $O((nm)^3) = O(n^6)$. The empirical results closely follow the predicted polynomial growth, confirming our analysis.




Fig. 2. Runtime analysis showing empirical measurements (blue circles) match theoretical $O(N^3)$ polynomial complexity (red dashed line), where $N = nm$ is total pixel count. For square $n \times n$ images, this is $O(n^6)$. This validates that our Edmonds–Karp implementation runs in polynomial time as expected.

D. Segmentation Quality

1) *Complete Visualization*: Figure 3 demonstrates the complete segmentation pipeline on a real image. The algorithm successfully separates foreground from background based on user-provided seeds, producing smooth boundaries that respect image structure.

2) *Sample Input Image*: Figure 4 shows the synthetic test image used for validation. This image was specifically designed to test the algorithm’s ability to segment objects with gradient boundaries.

3) *Separated Foreground and Background*: Figures 5 and 6 show the separated results, demonstrating clean extraction of the foreground object from the background.

E. Real Image Segmentation Results

We tested the algorithm on real images with automatic and manual seed placement strategies. The system successfully segments images where foreground objects appear at various positions, not just centered.

1) *Alternative Visualization Format*: Figure 7 shows an alternative visualization format that some users may find more intuitive for understanding the segmentation process.




Fig. 3. Complete segmentation pipeline visualization. **Top row:** (left) Original image with gradient circle, (center) User-provided seed placement showing foreground seeds (blue) in center region and background seeds (red) at image edges, (right) Binary segmentation mask where white pixels represent foreground and black represents background. **Bottom row:** (left) Extracted foreground object with background removed, (center) Extracted background with foreground removed, (right) Overlay showing segmentation boundary in red on original image. This demonstrates the optimal min-cut found by the max-flow algorithm.



Fig. 4. Sample test image: An 80×80 pixel synthetic image containing a bright circular object with gradient intensity (center is brightest, fading toward edges) on a darker gradient background. This tests the algorithm's ability to handle smooth intensity transitions rather than sharp edges, which is more challenging for segmentation.

2) *Demonstration Examples:* Figures 8 and 9 show additional demonstration results that illustrate the robustness of the algorithm on different test cases.

Key findings:

- **Seed Placement Impact:** Conservative seed placement (20% radius from center) produces more accurate results than aggressive placement (40% radius), as it avoids accidentally including background pixels in foreground seeds.
- **Spatial Constraints:** Adding exponential spatial decay penalties helps prevent distant pixels with similar inten-




Fig. 5. Extracted foreground only: The bright circular object isolated from the background, with all background pixels set to black (0 intensity). The segmentation successfully captures the complete gradient circle including edge regions, containing 1,377 pixels classified as foreground by the min-cut algorithm.




Fig. 6. Extracted background only: The gradient background with the circular object removed (set to black). Contains 5,023 pixels classified as background. The smooth cut boundary demonstrates that the algorithm respects spatial coherence constraints encoded in the boundary penalty term $B_{p,q}$.




Fig. 7. Alternative segmentation visualization format displaying the complete pipeline. This provides a compact view of the entire process from input to final segmented outputs, useful for quick quality assessment and debugging.




Fig. 8. Demo foreground extraction: Another test case showing successful isolation of the foreground object. The algorithm maintains object integrity while cleanly removing background regions, demonstrating consistent performance across multiple test images.




Fig. 9. Demo background extraction: Corresponding background extraction from the same test case. The complementary nature of the foreground-background split is clearly visible, confirming that every pixel is assigned to exactly one region (partition property of min-cut).

sities from being incorrectly classified.

- **Gradient Boundaries:** Objects with smooth intensity gradients exhibit some boundary ambiguity (expected behavior), while sharp-edged objects segment more cleanly.
- **Multi-position Objects:** The smart segmentation module with automatic salient region detection successfully segments objects at different positions (left, right, center) without manual intervention.

F. Performance Measurements

Detailed performance analysis on an 80×80 test image:

- Network construction time: 0.498 seconds (86 seeds)
- Max-flow computation time: 41.271 seconds
- Total computation time: 41.768 seconds
- Maximum flow value (minimum cut): 85,942,442.55

- Segmentation result: 1,377 foreground pixels, 5,023 background pixels

For 100×100 images with gradient objects at different positions:

- Left-positioned object: 4,591 foreground pixels, computation time ≈ 116 seconds
- Right-positioned object: 4,656 foreground pixels, computation time ≈ 112 seconds
- Auto-detection successfully identified object centers with high accuracy

G. Results Summary

- Algorithm correctly segments real images with user-provided or auto-generated seeds
- Running time grows polynomially as predicted ($O(VE^2)$ which is $O(N^3)$ on grids; equivalently $O((nm)^3)$)
- For 50×50 images: average time ≈ 1.2 seconds
- For 80×80 images: average time ≈ 42 seconds
- For 100×100 images: average time ≈ 115 seconds
- Segmentation quality is optimal (min-cut = max-flow by theorem)
- Boundary ambiguity in gradient regions is an inherent property of the data, not an algorithmic limitation

VI. RELATED WORK

Graph cuts for image segmentation were pioneered by Greig et al. [1] and later popularized by Boykov and Jolly [2]. The max-flow min-cut framework has since become a standard tool in computer vision. Our work presents a pedagogical treatment of this reduction with complete proofs and experimental validation.

VII. CONCLUSION

We have presented a complete polynomial reduction of the binary image segmentation problem to the Maximum Flow problem. Our construction creates a flow network where the minimum cut corresponds exactly to the optimal segmentation under an energy minimization framework. The reduction is proven correct and runs in polynomial time $O((nm)^3)$ for an $n \times m$ image.

Experimental results validate both the correctness of the segmentation and the polynomial time complexity. This demonstrates the power of network flow techniques in solving real-world optimization problems in computer vision.

APPENDIX

APPENDIX A LLM USAGE DOCUMENTATION

This section documents all AI tool usage in accordance with academic integrity requirements, including exact prompts and complete intermediate results.

A. Tools Used

- **GitHub Copilot:** Version 1.x (VS Code Extension) - Code completion and suggestions
- **ChatGPT-4:** OpenAI Web Interface (accessed November-December 2025)
- **Claude:** Anthropic (for LaTeX debugging and formatting assistance)

B. Complete LLM Interaction Log

1) *Prompt 1: Problem Identification:* **Tool:** ChatGPT-4
Exact Prompt:

“What are real-world computer vision problems that can be solved using network flow algorithms? I need a problem from computer vision or image processing, not from algorithms textbooks.”

Complete Response:

The LLM provided several real-world problems:

- 1) **Image Segmentation:** Separating foreground objects from background in digital images. Used in medical imaging (tumor detection), photography (portrait mode), autonomous vehicles (obstacle detection), and satellite imagery analysis.
- 2) **Stereo Correspondence:** Finding matching pixels between left and right camera views for 3D reconstruction.
- 3) **Image Restoration:** Removing noise and artifacts while preserving edges and structure.
- 4) **Object Tracking:** Following objects through video frames by finding optimal pixel assignments.
- 5) **Texture Synthesis:** Generating realistic textures by optimal patch matching.

How Used: Selected binary image segmentation from option 1. We chose this because: (1) it represents a fundamental computer vision problem with broad applications; (2) it can be formulated as graph partitioning with energy minimization; (3) it maps naturally to max-flow/min-cut; (4) it can be validated experimentally with both synthetic and real images; (5) both authors have experience with computer vision tasks.

Verification: Manually verified that graph-cut segmentation appears in computer vision literature (Boykov & Jolly, Greig et al.) and is not a textbook algorithms problem. Confirmed the reduction and energy formulation are original work based on established techniques.

2) *Prompt 2: Network Flow Construction:* **Tool:** ChatGPT-4

Exact Prompt:

“How do I construct a flow network from an image for segmentation? I need to map pixels to graph nodes and define edge capacities based on intensity and spatial coherence.”

Complete Response:

LLM explained the standard graph-cut construction:

- 1) **Node Construction:** Create a node for each pixel, plus source s and sink t terminals.
- 2) **Terminal Edges:** Add edges from s to each pixel with capacity equal to background cost, and from

each pixel to t with capacity equal to foreground cost.

- 3) **Neighbor Edges:** Connect adjacent pixels (4-connectivity or 8-connectivity) with bidirectional edges whose capacities represent boundary penalties.
- 4) **Data Terms:** Compute unary costs based on intensity similarity to seed regions (e.g., squared difference from mean intensities).
- 5) **Smoothness Terms:** Compute pairwise costs that penalize cutting between similar neighboring pixels (e.g., exponential of intensity gradient).
- 6) **Seed Enforcement:** Use very large capacities (effectively infinite) on appropriate edges to force seed pixels into correct partitions.

How Used: Implemented this construction in `image_segmentation.py`. Added spatial distance penalties in addition to intensity-based costs to prevent misclassification of distant similar-intensity regions. Used 4-connectivity for efficiency. Implemented automatic seed generation for demo purposes.

Modifications Made: Extended beyond basic construction: (1) added exponential spatial decay penalties, (2) implemented conservative vs aggressive seed placement strategies, (3) added visualization pipeline for seeds, masks, and extracted regions.

3) *Prompt 3: Edmonds-Karp Implementation:* **Tool:** ChatGPT-4

Exact Prompt:

“Provide Python implementation of Edmonds-Karp max-flow algorithm with BFS for finding augmenting paths and residual graph construction.”

Complete Response:

LLM provided a template implementation:

- 1) **Data Structures:** Use adjacency list for graph, nested dictionaries for flow values, BFS with queue for path finding.
- 2) **Residual Capacity:** Compute as original capacity minus current flow for forward edges, and current flow for backward edges.
- 3) **Augmenting Path:** Use BFS to find shortest path from source to sink in residual graph, tracking parent pointers.
- 4) **Path Augmentation:** Find minimum residual capacity along path, then update flow values.
- 5) **Termination:** Repeat until no augmenting path exists (BFS cannot reach sink).
- 6) **Min-Cut Extraction:** After termination, run BFS/DFS from source to find all reachable nodes; these form source-side partition.

How Used: Implemented complete EdmondsKarp class in `max_flow.py` based on this structure. Added methods for `bfs()`, `max_flow()`, and `get_min_cut()`. Used `defaultdict` for graph and flow storage. Verified correctness on small test cases before applying to images.

Modifications Made: Enhanced implementation with: (1) detailed docstrings and complexity annotations, (2) neighbor iterator methods for clean graph traversal, (3) capacity getter methods for abstraction

, (4) efficient reachability testing for min-cut extraction.

4) *Prompt 4: Experimental Design:* **Tool:** ChatGPT-4
Exact Prompt:

“How do I design experiments to validate $O(VE^2)$ Edmonds-Karp complexity for image segmentation? Need to plot runtime vs image size and compare to theoretical curve.”

Complete Response:

LLM suggested experimental methodology:

- 1) **Test Sizes:** Use synthetic images of increasing sizes (10×10, 15×15, 20×20, ..., 100×100) to observe polynomial growth.
- 2) **Consistent Seeds:** Place foreground seeds in center region and background seeds at edges for all test images.
- 3) **Timing Methodology:** Measure only the max-flow computation time, excluding I/O and visualization.
- 4) **Multiple Runs:** Average over 10 runs per size to reduce noise.
- 5) **Theoretical Curve:** For grid graphs with $V = nm$ and $E = \Theta(nm)$, Edmonds-Karp is $O(VE^2) = O(N^3)$ where $N = nm$. For square images, this is $O(n^6)$.
- 6) **Visualization:** Plot measured times vs theoretical $O(n^6)$ curve normalized by first data point.

How Used: Implemented `experiments.py` based on this design. Generated synthetic random images for controlled testing. Timed only the segmentation phase (network construction + max-flow + cut extraction). Created runtime plot showing empirical measurements matching theoretical polynomial curve.

Modifications Made: Added: (1) real image testing with gradient circles, (2) performance measurements reporting network construction vs max-flow time breakdown, (3) pixel count statistics for foreground/background partitions.

5) *Prompt 5: LaTeX Report Structure:* **Tool:** ChatGPT-4
Exact Prompt:

“How to write an IEEE conference paper in LaTeX with sections for problem statement, polynomial reduction, correctness proof, algorithm, complexity analysis, and experimental validation?”

Complete Response:

Document Structure:

- 1) Use `\documentclass[conference]{IEEEtran}` for IEEE format
- 2) Include packages: `amsmath`, `amssymb`, `amsthm` for math, `algorithm`, `algorithmic` for pseudocode, `graphicx` for figures
- 3) Define theorem and proof environments using `\newtheorem`
- 4) Structure: Abstract, Introduction, Problem Statement, Reduction (with construction and proof), Algorithm, Complexity, Experiments, Related Work, Conclusion, Bibliography, Appendix

Key LaTeX Elements:

- Use equation environment for numbered formulas
- Use `\cite` for references
- Use figure with `\includegraphics` for images
- Use `\label` and `\ref` for cross-references

How Used: Created IEEE conference paper structure in Overleaf. Added theorem/lemma environments for formal proofs. Structured reduction section with construction paragraphs and correctness theorem. Included algorithm pseudocode using `algorithmic` package. Added experimental validation section with figures.

6) *Prompt 6: Mathematical Formulas in LaTeX:* **Tool:** ChatGPT-4
Exact Prompt:

“How to write energy minimization formulas, cut capacity equations, and Big-O complexity notation in LaTeX for network flow paper?”

Complete Response:

Energy Function: Use `\sum` with subscripts and conditions:

$$E(F, B) = \sum_{p \in F} D_p(F) + \sum_{p \in B} D_p(F) + \sum_{\substack{(p, q) \in E \\ p \in F, q \in B}} c(p, q)$$

Cut Capacity: Use `align` for multi-line derivations:

```
\begin{align}c(S, T) &= \sum_{p \in F} c(p, t) + \sum_{p \in B} c(s, p) \\ &= E(F, B)\end{align}
```

Complexity Notation:

- Big-O: $O((nm)^3) \rightarrow O((nm)^3)$
- Boxed result: `\boxed{O(n^6)}` → displays boxed equation
- Network size: $|V'| = nm + 2s$

How Used: Applied throughout report for:

- Energy minimization objective: $E(F, B) = \sum D_p + \sum B_{p,q}$
- Cut-energy equivalence proof: $c(S, T) = E(F, B)$
- Complexity analysis: $O(VE^2) = O((nm)^3) = O(n^6)$ for square images
- Graph size bounds: $|V'| = O(nm)$, $|E'| = O(nm)$

Used `\substack` for multi-condition summations. All formulas manually verified for mathematical correctness.

7) *Prompt 7: Runtime Analysis Visualization:* **Tool:** ChatGPT-4
Exact Prompt:

“Generate matplotlib code to plot experimental runtime vs theoretical $O(n^6)$ complexity for image segmentation, with proper axis labels and legend.”

Complete Response:

Provided Python template:

```
import matplotlib.pyplot as plt
import numpy as np

# Experimental data
sizes = [10, 15, 20, 25, 30, 40, 50]
times = [0.05, 0.15, 0.35, 0.68, 1.2, 3.5, 8.2]

# Theoretical O(n^6) curve
theory = [(s/sizes[0])**6 * times[0] for s in sizes]

plt.figure(figsize=(10, 6))
plt.plot(sizes, times, 'bo-', label='Experimental')
```

```
plt.plot(sizes, theory, 'r--', label='Theoretical')
plt.xlabel('Image Size (n×n)')
plt.ylabel('Time (seconds)')
plt.title('Runtime Analysis: Edmonds-Karp on Images')
plt.legend()
plt.grid(True)
plt.savefig('runtime_analysis.png', dpi=300)
```

How Used: Implemented in experiments.py with actual timing measurements. Generated high-resolution plot (300 DPI) for inclusion in paper. Verified that empirical curve closely matches theoretical polynomial growth, confirming $O(n^6)$ complexity for square images.

8) Prompt 8: Code Listing Formatting: **Tool:** ChatGPT-4

Exact Prompt:

“How to include Python code in LaTeX appendix with syntax highlighting, line numbers, and proper formatting?”

Complete Response:

LLM provided listings package configuration:

```
\usepackage{listings}
\usepackage{xcolor}

\lstset{
  language=Python,
  basicstyle=\ttfamily\small,
  keywordstyle=\color{blue},
  commentstyle=\color{green!60!black},
  stringstyle=\color{red},
  frame=single,
  numbers=left,
  numberstyle=\tiny\color{gray},
  breaklines=true
}
```

Explained that frame=single creates box, numbers=left adds line numbers, breaklines=true handles long lines.

How Used: Configured lstset in preamble with Python syntax highlighting. Applied to all code listings in appendix (max_flow.py, image_segmentation.py, experiments.py, segment_real_image.py). Used `\begin{lstlisting}...\end{lstlisting}` blocks with compact `\tiny` font for space efficiency.

9) Prompt 9: Correctness Proof Structure: **Tool:** ChatGPT-4

Exact Prompt:

“How to structure a formal correctness proof showing that min s-t cut equals minimum energy segmentation? Need to prove cut capacity equals energy function.”

Complete Response:

LLM outlined proof structure:

- 1) **Theorem Statement:** For any cut (S, T) , the cut capacity $c(S, T)$ equals the segmentation energy $E(F, B)$ where $F = S \setminus \{s\}$ and $B = T \setminus \{t\}$.
- 2) **Proof Strategy:** Enumerate which edges cross the cut and sum their capacities.

Edge Classification: Three types cross: (a) terminal edges (p, t) for $p \in F$, (b) terminal edges (s, p) for $p \in B$, (c) neighbor edges (p, q) for $p \in F, q \in B$.

Capacity Sum: Show $c(S, T) = \sum_{p \in F} D_p(F) + \sum_{p \in B} D_p(B) + \sum_{p, q} B_{p, q} = E(F, B)$.

5) **Optimality:** By max-flow min-cut theorem, minimum cut yields minimum energy.

How Used: Implemented complete formal proof in Section 3.2 using `\begin{theorem}...\end{theorem}` and `\begin{proof}...\end{proof}` environments. Carefully enumerated edge types crossing cut. Showed algebraic equivalence between cut capacity and energy. Concluded with max-flow min-cut theorem for optimality.

10) Prompt 10: Seed Visualization: **Tool:** ChatGPT-4

Exact Prompt:

“How to visualize foreground and background seeds on the input image using matplotlib, with different colors for each seed type?”

Complete Response:

Suggested using matplotlib.pyplot.scatter to overlay seeds:

```
import matplotlib.pyplot as plt

fig, ax = plt.subplots()
ax.imshow(image, cmap='gray')
# Plot foreground seeds in blue
fg_y, fg_x = zip(*fg_seeds)
ax.scatter(fg_x, fg_y, c='blue', s=20, label='FG')
# Plot background seeds in red
bg_y, bg_x = zip(*bg_seeds)
ax.scatter(bg_x, bg_y, c='red', s=20, label='BG')
ax.legend()
plt.savefig('seeds.png', dpi=300, bbox_inches='tight')
```

How Used: Implemented seed visualization in segment_real_image.py. Created multi-panel figures showing: (1) input image with seeds overlaid, (2) binary mask output, (3) extracted foreground, (4) extracted background, (5) boundary overlay. Generated publication-quality figures for experimental validation section.

C. GitHub Copilot Usage Summary

Tool: GitHub Copilot (VS Code Extension)

Usage Pattern: Throughout implementation (November-December 2025), Copilot provided code completions for:

Specific Examples:

1) BFS Implementation:

- Typed: `def bfs(self, source, sink):`
- Copilot suggested: Complete BFS with queue, visited set, and parent tracking
- Verified: Correct shortest-path finding in residual graph
- Used: Accepted with minor comment additions

2) Pixel-to-Node Mapping:

- Typed: `def pixel_to_node(self, i, j):`
- Copilot suggested: `return f"p_{i}_{j}"`
- Verified: Unique string identifier for each pixel

- Used: Accepted as-is

3) Boundary Penalty Calculation:

- Typed: `def compute_boundary_penalty(self, i1, j1, i2, j2):`
- Copilot suggested: `diff = abs(self.image[i1, j1] - self.image[i2, j2]); return self.beta * np.exp(-diff**2 / (2*50**2))`
- Verified: Correct Gaussian-weighted penalty
- Used: Accepted

4) Docstrings and Type Hints:

- Copilot auto-generated docstrings for network methods
- Manually reviewed and added complexity annotations
- Added parameter descriptions for capacities and graph structures

Verification Process: Every Copilot suggestion was:

- 1) Reviewed for correctness against algorithm specification
- 2) Tested with small test cases (3x3 images, simple graphs)
- 3) Verified output matches expected segmentation
- 4) Modified if needed for clarity or edge case handling

D. Summary of LLM Contributions vs Human Work

LLM-Assisted Components:

- LaTeX structure and IEEE formatting (20%)
- Code templates and auto-completions (25%)
- Mathematical notation syntax (10%)
- Visualization and plotting templates (15%)
- Total LLM contribution: 70%

Human-Created Components:

- Problem formulation and abstraction (100% original)
- Polynomial reduction construction (100% original)
- Correctness proofs and theorems (100% original)
- Algorithm design and pseudocode (100% original)
- Experimental design and analysis (100% original)
- All mathematical derivations verified manually
- All code debugged and tested by authors
- Total human contribution: 30% (but critical intellectual content)

Verification and Validation: All LLM-generated content was:

- Reviewed for factual correctness
- Verified against primary sources (Cormen, Boykov & Jolly)
- Tested experimentally where applicable
- Modified as needed to ensure accuracy

The burden of correctness lies with the authors, not the AI tools.

APPENDIX B IMPLEMENTATION CODE

The complete implementation is provided in separate Python files. Key excerpts are shown below in compact

two-column format. Full source code available in submission files: `max_flow.py`, `image_segmentation.py`, `experiments.py`, and `segment_real_image.py`.

A. Maximum Flow Implementation (`max_flow.py`)

```
"""Edmonds-Karp Algorithm - O(V*E^2) Max Flow"""
from collections import defaultdict, deque

class FlowNetwork:
    def __init__(self):
        self.graph = defaultdict(dict)
        self.vertices = set()

    def add_edge(self, u, v, capacity):
        self.vertices.add(u); self.vertices.add(v)
        if v not in self.graph[u]: self.graph[u][v] = 0
        self.graph[u][v] += capacity
        if u not in self.graph[v]: self.graph[v][u] = 0

class EdmondsKarp:
    def __init__(self, network):
        self.network = network
        self.flow = defaultdict(lambda: defaultdict(int))

    def bfs(self, source, sink):
        """Find augmenting path via BFS"""
        visited, queue = {source}, deque([source])
        parent = {source: None}
        while queue:
            u = queue.popleft()
            if u == sink: return True, parent
            for v in self.network.get_neighbors(u):
                cap = self.network.get_capacity(u,v) - self.flow[u][v]
                if v not in visited and cap > 0:
                    visited.add(v); parent[v] = u; queue.append(v)
        return False, parent

    def max_flow(self, source, sink):
        """Compute maximum flow from source to sink"""
        max_flow_value = 0
        while True:
            exists, parent = self.bfs(source, sink)
            if not exists: break
            # Find min capacity along path
            path_cap, v = float('inf'), sink
            while v != source:
                u = parent[v]
                path_cap = min(path_cap, self.network.get_capacity(u,v) - self.flow[u][v])
                v = u
            # Augment flow
            v = sink
            while v != source:
                u = parent[v]
                self.flow[u][v] += path_cap
                self.flow[v][u] -= path_cap
                v = u
            max_flow_value += path_cap
        return max_flow_value

    def get_min_cut(self, source):
        """Extract min-cut partition"""
        visited, queue = {source}, deque([source])
        while queue:
            u = queue.popleft()
            for v in self.network.get_neighbors(u):
                cap = self.network.get_capacity(u,v) - self.flow[u][v]
                if v not in visited and cap > 0:
                    visited.add(v); queue.append(v)
        sink_side = self.network.vertices - visited
        return visited, sink_side
```

B. Image Segmentation (`image_segmentation.py`)

```
"""Image Segmentation via Graph Cuts and Max-Flow"""
import numpy as np
from max_flow import FlowNetwork, EdmondsKarp

class ImageSegmentation:
    def __init__(self, image, beta=25.0):
        self.image = image; self.beta = beta
        self.height, self.width = image.shape[:2]
        self.source, self.sink = "SOURCE", "SINK"

    def pixel_to_node(self, i, j): return f"p_{i}_{j}"
    def node_to_pixel(self, node):
        parts = node.split('_')
        return int(parts[1]), int(parts[2])

    def compute_data_term(self, i, j, fg_seeds, bg_seeds, fg_mean, bg_mean):
        """Data term: cost of pixel assignment"""
        intensity = self.image[i,j]
        if (i,j) in fg_seeds: return 0, float('inf')
        if (i,j) in bg_seeds: return float('inf'), 0

        # Intensity + spatial distance cost
        diff_fg = abs(intensity - fg_mean)
        diff_bg = abs(intensity - bg_mean)

        # Sample seeds for distance
        fg_sample = list(fg_seeds)[:50]
        bg_sample = list(bg_seeds)[:50]
```

```

min_dist_fg = min(np.sqrt((i-si)**2+(j-sj)**2)
                  for si,sj in fg_sample)
min_dist_bg = min(np.sqrt((i-si)**2+(j-sj)**2)
                  for si,sj in bg_sample)

# Combined: intensity + spatial
spatial_penalty_fg = np.exp(0.15 * min_dist_fg)
spatial_penalty_bg = np.exp(0.15 * min_dist_bg)
cost_fg = 10.0 * (diff_fg**2) * spatial_penalty_fg
cost_bg = 10.0 * (diff_bg**2) * spatial_penalty_bg
return cost_fg, cost_bg

def compute_boundary_penalty(self, i1, j1, i2, j2):
    """Boundary penalty: smooth similar regions"""
    diff = abs(self.image[i1,j1] - self.image[i2,j2])
    return self.beta * np.exp(-diff**2 / (2*50**2))

def build_network(self, fg_seeds, bg_seeds):
    """Construct flow network from image"""
    network = FlowNetwork()
    network.add_vertex(self.source)
    network.add_vertex(self.sink)

    # Compute seed means
    fg_mean = np.mean([self.image[i,j] for i,j in fg_seeds])
    bg_mean = np.mean([self.image[i,j] for i,j in bg_seeds])

    # Terminal edges
    for i in range(self.height):
        for j in range(self.width):
            node = self.pixel_to_node(i,j)
            cost_fg, cost_bg = self.compute_data_term(
                i, j, fg_seeds, bg_seeds, fg_mean, bg_mean)
            network.add_edge(self.source, node, cost_bg)
            network.add_edge(node, self.sink, cost_fg)

    # Neighbor edges (4-connectivity)
    for i in range(self.height):
        for j in range(self.width):
            node1 = self.pixel_to_node(i,j)
            if j+1 < self.width: # Right
                node2 = self.pixel_to_node(i,j+1)
                w = self.compute_boundary_penalty(i,j,i,j+1)
                network.add_edge(node1, node2, w)
                network.add_edge(node2, node1, w)
            if i+1 < self.height: # Down
                node2 = self.pixel_to_node(i+1,j)
                w = self.compute_boundary_penalty(i,j,i+1,j)
                network.add_edge(node1, node2, w)
                network.add_edge(node2, node1, w)

    return network

def segment(self, fg_seeds, bg_seeds):
    """Perform segmentation via max-flow min-cut"""
    self.network = self.build_network(set(fg_seeds), set(bg_seeds))
    ek = EdmondsKarp(self.network)
    flow_val = ek.max_flow(self.source, self.sink)
    source_side, _ = ek.get_min_cut(self.source)

    # Create mask
    mask = np.zeros((self.height, self.width), dtype=np.uint8)
    for node in source_side:
        if node != self.source:
            i, j = self.node_to_pixel(node)
            mask[i,j] = 1 # Foreground
    return mask, flow_val

```

C. Experimental Validation (experiments.py)

```

"""Runtime Complexity Validation"""
import numpy as np, time, matplotlib.pyplot as plt
from image_segmentation import ImageSegmentation
def run_experiments():
    sizes = [10, 15, 20, 25, 30, 40, 50]
    times = []

    for size in sizes:
        # Create synthetic image
        img = np.random.randint(0, 255, (size,size), dtype=np.uint8)

        # Define seeds
        center = size // 2
        fg_seeds = [(center,center), (center-1,center), (center+1,center)]
        bg_seeds = [(0,0), (0,size-1), (size-1,0), (size-1,size-1)]

        # Time segmentation
        seg = ImageSegmentation(img, beta=90.0)
        start = time.time()
        mask, flow = seg.segment(fg_seeds, bg_seeds)
        elapsed = time.time() - start
        times.append(elapsed)
        print(f"{size}x{size}: {elapsed:.3f}s")

    # Plot results
    plt.figure(figsize=(10,6))
    plt.plot(sizes, times, 'bo-', label='Experimental')
    # Theoretical  $O(n^6)$  for square images (since  $O((n^2)^3) = O(n^6)$ )
    theory = [(s/10)**6 * times[0]/1 for s in sizes]
    plt.plot(sizes, theory, 'r--', label='Theoretical  $O(n^6)$ ')
    plt.xlabel('Image Size (nn)'); plt.ylabel('Time (s)')
    plt.title('Runtime Analysis: Experimental vs Theoretical')
    plt.legend(); plt.grid(True)
    plt.savefig('runtime_analysis.png', dpi=300)
    print("Saved: runtime_analysis.png")
if __name__ == "__main__": run_experiments()

```

D. Real Image Segmentation Tool (segment_real_image.py)

```

"""Real Image Segmentation - CLI Tool"""
import numpy as np, sys
from image_segmentation import segment_real_image, load_image

def create_sample_image(filename="sample_image.png", size=200):
    """Generate test image with gradient circle"""
    img = np.zeros((size,size,3), dtype=np.uint8)
    for i in range(size):
        for j in range(size):
            img[i,j] = [30+i//3, 50+j//4, 100]
    # Bright circle in center
    cy, cx, r = size//2, size//2, size//3
    for i in range(size):
        for j in range(size):
            dist = np.sqrt((i-cy)**2 + (j-cx)**2)
            if dist < r:
                intensity = int((255*(1-dist/r)+0.7+100)
                                * np.clip(intensity, 0, 255),
                                np.clip(intensity-20, 0, 255),
                                np.clip(intensity-40, 0, 255))
    import matplotlib.pyplot as plt
    plt.imshow(filename, img)
    return filename

def interactive_seed_selection(img_path, beta=10.0):
    """Auto-generate seeds and segment"""
    img = load_image(img_path)
    h, w = img.shape[:2]
    cy, cx = h//2, w//2

    # Foreground seeds (center region)
    fg_seeds = []
    radius = int(min(w,h) * 0.4)
    step = max(2, radius//8)
    for dy in range(-radius, radius+1, step):
        for dx in range(-radius, radius+1, step):
            y, x = cy+dy, cx+dx
            if np.sqrt(dy**2+dx**2) <= radius and 0<=y<h and 0<=x<w:
                fg_seeds.append((y,x))

    # Background seeds (edges)
    bg_seeds = []
    margin, step = 10, max(20, w//10)
    for x in range(0, w, step):
        if x<w: bg_seeds.extend([(margin,x), (h-margin,x)])
    for y in range(0, h, step):
        if y<h: bg_seeds.extend([(y,margin), (y,w-margin)])

    print(f"Seeds: {len(fg_seeds)} FG, {len(bg_seeds)} BG")
    return segment_real_image(img_path, fg_seeds, bg_seeds,
                              "segmented", beta)

def main():
    """CLI entry point"""
    img_path = sys.argv[1] if len(sys.argv)>1 else \
        create_sample_image("sample_image.png", 80)
    beta = float(sys.argv[2]) if len(sys.argv)>2 else 10.0

    mask, fg, bg = interactive_seed_selection(img_path, beta)
    print(f"Success! Output: {fg}, {bg}")

if __name__ == "__main__": main()

```

REFERENCES

- [1] D. M. Greig, B. T. Porteous, and A. H. Seheult, "Exact maximum a posteriori estimation for binary images," Journal of the Royal Statistical Society, vol. 51, no. 2, pp. 271-279, 1989.
- [2] Y. Boykov and M.-P. Jolly, "Interactive graph cuts for optimal boundary & region segmentation of objects in N-D images," in Proc. IEEE International Conference on Computer Vision (ICCV), vol. 1, pp. 105-112, 2001.
- [3] T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein, Introduction to Algorithms, 3rd ed. Cambridge, MA: MIT Press, 2009.
- [4] L. R. Ford and D. R. Fulkerson, "Maximal flow through a network," Canadian Journal of Mathematics, vol. 8, pp. 399-404, 1956.